

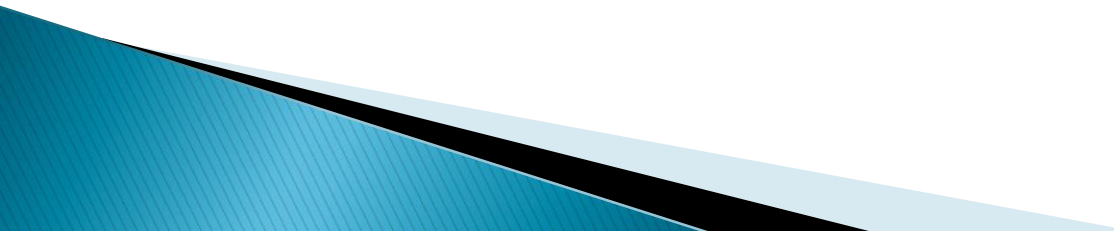
Apache Kafka Development

Rajesh Pasham

Consumer Group Example

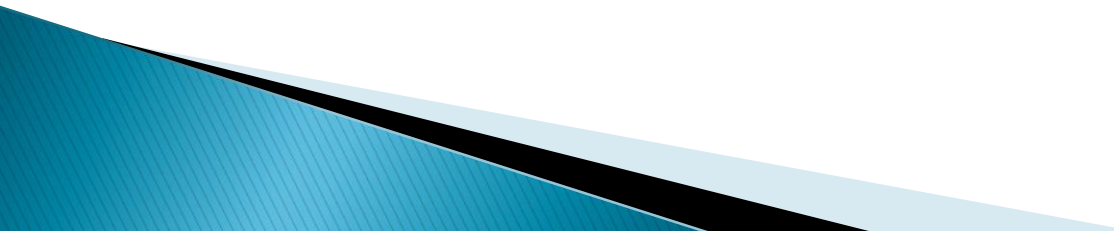
- ▶ Consumer group is a multi-threaded or multi-machine consumption from Kafka topics.

Consumer Group

- ▶ Consumers can join a group by using the `samegroup.id`.
 - ▶ The maximum parallelism of a group is that the number of consumers in the group \leftarrow no of partitions.
 - ▶ Kafka assigns the partitions of a topic to the consumer in a group, so that each partition is consumed by exactly one consumer in the group.
 - ▶ Kafka guarantees that a message is only ever read by a single consumer in the group.
 - ▶ Consumers can see the message in the order they were stored in the log.
- 

Consumer Group Example

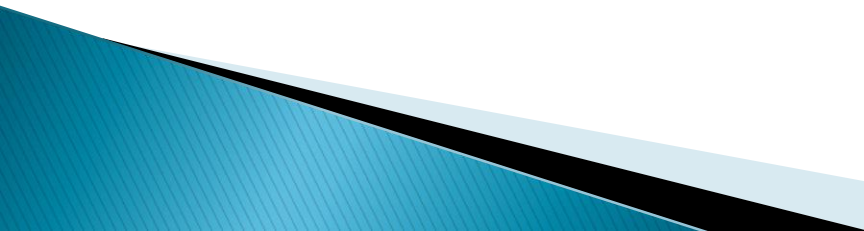
Re-balancing of a Consumer

- ▶ Adding more processes/threads will cause Kafka to re-balance.
 - ▶ If any consumer or broker fails to send heartbeat to ZooKeeper, then it can be re-configured via the Kafka cluster.
 - ▶ During this re-balance, Kafka will assign available partitions to the available threads, possibly moving a partition to another process.
- 

Consumer Group Example

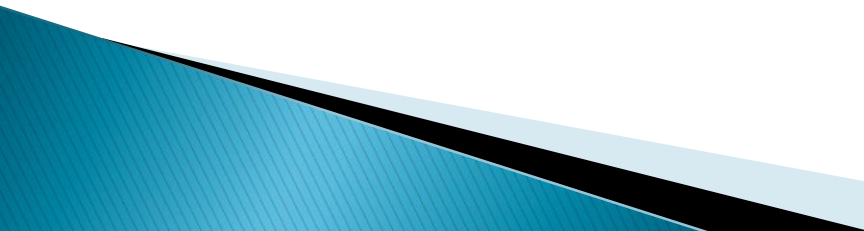
```
import java.util.Properties;
import java.util.Arrays;
import org.apache.kafka.clients.consumer.KafkaConsumer;
import org.apache.kafka.clients.consumer.ConsumerRecords;
import org.apache.kafka.clients.consumer.ConsumerRecord;

public class ConsumerGroup {
    public static void main(String[] args) throws Exception {
        if(args.length < 2){
            System.out.println("Usage: consumer <topic> <groupname>");
            return;
        }
    }
}
```



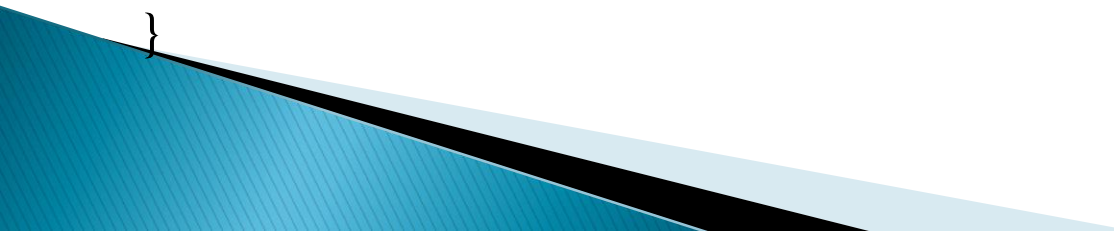
Consumer Group Example

```
String topic = args[0].toString();
String group = args[1].toString();
Properties props = new Properties();
props.put("bootstrap.servers", "localhost:9092");
props.put("group.id", group);
props.put("enable.auto.commit", "true");
props.put("auto.commit.interval.ms", "1000");
props.put("session.timeout.ms", "30000");
props.put("key.deserializer",
    "org.apache.kafka.common.serialization.StringDeserializer");
props.put("value.deserializer",
    "org.apache.kafka.common.serialization.StringDeserializer");
```



Consumer Group Example

```
KafkaConsumer<String, String> consumer = new
KafkaConsumer<String, String>(props);
consumer.subscribe(Arrays.asList(topic));
System.out.println("Subscribed to topic " + topic);
int i = 0;
while (true) {
    ConsumerRecords<String, String> records = consumer.poll(100);
    for (ConsumerRecord<String, String> record : records)
        System.out.printf("offset = %d, key = %s, value = %s\n",
            record.offset(), record.key(), record.value());
    }
}
```



Consumer Group Example

Compilation

- ▶ `javac -cp "/path/to/kafka/kafka_2.11-0.9.0.0/libs/*" ConsumerGroup.java`

Execution

- ▶ `>>java -cp "/path/to/kafka/kafka_2.11-0.9.0.0/libs/*":.
ConsumerGroup <topic-name> my-group`
- ▶ `>>java -cp "/path/to/kafka/kafka_2.11-0.9.0.0/libs/*":.
ConsumerGroup <topic-name> my-group`

Here we have created a sample group name as my-group with two consumers. Similarly, you can create your group and number of consumers in the group.

Consumer Group Example

Input

- ▶ Open producer CLI and send some messages like
 - Test consumer group 01
 - Test consumer group 02

Output of the First Process

Subscribed to topic Hello-kafka

offset = 3, key = null, value = Test consumer group 01

Output of the Second Process

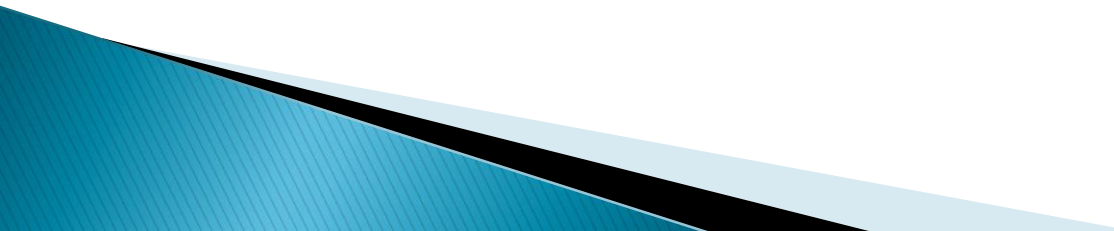
Subscribed to topic Hello-kafka

offset = 3, key = null, value = Test consumer group 02

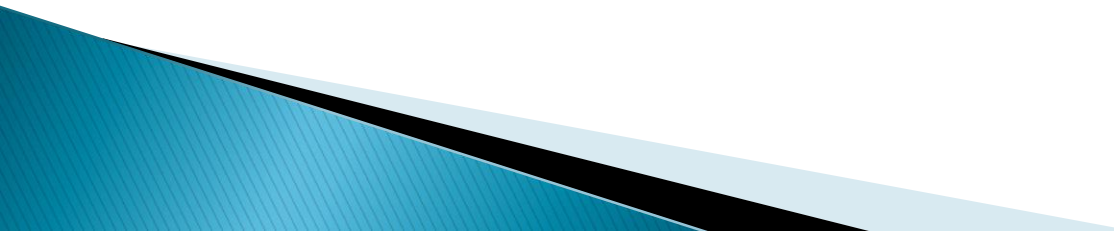


Kafka Connectors

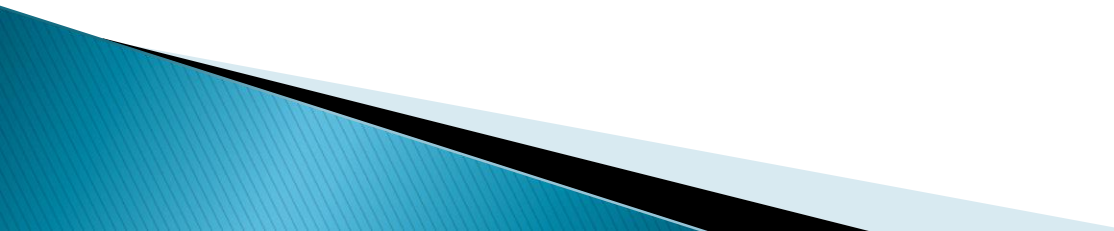
Overview

- ▶ Apache Kafka is a distributed streaming platform.
 - ▶ In this topic, we'll learn how to use Kafka Connectors.
 - ▶ We'll have a look at:
 - Different types of Kafka Connectors
 - Features and modes of Kafka Connect
 - Connectors configuration using property files as well as the REST API
- 


Basics of Kafka Connect and Kafka Connectors

- ▶ Kafka Connect is a framework for connecting Kafka with external systems such as databases, key-value stores, search indexes, and file systems, using so-called *Connectors*.
 - ▶ Kafka Connectors are ready-to-use components, which can help us to import data from external systems into Kafka topics and export data from Kafka topics into external systems.
 - ▶ We can use existing connector implementations for common data sources and sinks or implement our own connectors.
 - ▶ A *source connector* collects data from a system.
 - ▶ Source systems can be entire databases, streams tables, or message brokers.
 - ▶ A source connector could also collect metrics from application servers into Kafka topics, making the data available for stream processing with low latency.
- 

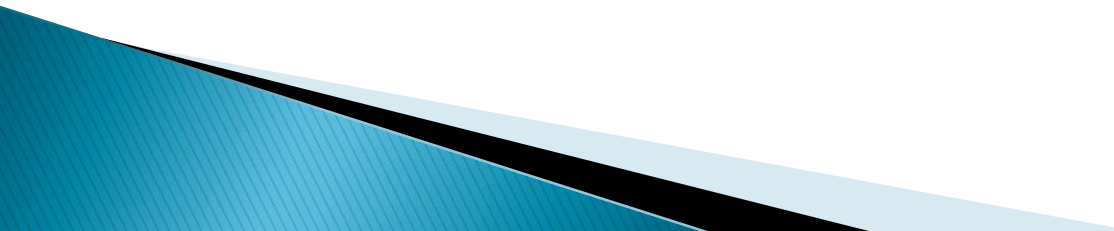
Basics of Kafka Connect and Kafka Connectors

- ▶ A *sink connector* delivers data from Kafka topics into other systems, which might be indexes such as Elasticsearch, batch systems such as Hadoop, or any kind of database.
 - ▶ Some connectors are maintained by the community, while others are supported by Confluent or its partners.
 - ▶ Really, we can find connectors for most popular systems, like S3, JDBC, and Cassandra, just to name a few.
- 

Features

- ▶ Kafka Connect features include:
 - A framework for connecting external systems with Kafka – it simplifies the development, deployment, and management of connectors
 - Distributed and standalone modes – it helps us to deploy large clusters by leveraging the distributed nature of Kafka, as well as setups for development, testing, and small production deployments
 - REST interface – we can manage connectors using a REST API
 - Automatic offset management – Kafka Connect helps us to handle the offset commit process, which saves us the trouble of implementing this error-prone part of connector development manually
- 

Features

- ▶ Kafka Connect features include:
 - Distributed and scalable by default – Kafka Connect uses the existing group management protocol; we can add more workers to scale up a Kafka Connect cluster
 - Streaming and batch integration – Kafka Connect is an ideal solution for bridging streaming and batch data systems in connection with Kafka's existing capabilities
 - Transformations – these enable us to make simple and lightweight modifications to individual messages
- 

Quick Start Kafka Connect

- ▶ For starters, we'll discuss the principle of Kafka Connect, using its most basic Connectors, which are the file *source* connector and the file *sink* connector.
- ▶ Conveniently, kafka comes with both of these connectors, as well as reference configurations.

Quick Start Kafka Connect

Source Connector Configuration

- ▶ For the source connector, the reference configuration is available at `$KAFKA_HOME/config/connect-file-source.properties`:
 - `name=local-file-source`
 - `connector.class=FileStreamSource`
 - `tasks.max=1`
 - `topic=connect-test`
 - `file=test.txt`

Quick Start Kafka Connect

Source Connector Configuration

- ▶ This configuration has some properties that are common for all source connectors:
 - *name* is a user-specified name for the connector instance
 - *connector.class* specifies the implementing class, basically the kind of connector
 - *tasks.max* specifies how many instances of our source connector should run in parallel, and
 - *topic* defines the topic to which the connector should send the output
- ▶ In this case, we also have a connector-specific attribute:
 - *File* defines the file from which the connector should read the input

Quick Start Kafka Connect

Source Connector Configuration

- ▶ For this to work then, let's create a basic file with some content:
 - `echo -e "foo\nbar\n" > $KAFKA_HOME/test.txt`
- ▶ Note that the working directory is `$KAFKA_HOME`.

Quick Start Kafka Connect

Sink Connector Configuration

- ▶ For our sink connector, we'll use the reference configuration at *\$KAFKA_HOME/config/connect-file-sink.properties*:
 - `name=local-file-sink`
 - `connector.class=FileStreamSink`
 - `tasks.max=1`
 - `file=test.sink.txt`
 - `topics=connect-test`
- ▶ Logically, it contains exactly the same parameters, though this time *connector.class* specifies the sink connector implementation, and *file* is the location where the connector should write the content.

Quick Start Kafka Connect

Worker Config

- ▶ Finally, we have to configure the Connect worker, which will integrate our two connectors and do the work of reading from the source connector and writing to the sink connector.
- ▶ For that, we can use *\$KAFKA_HOME/config/connect-standalone.properties*:
 - bootstrap.servers=localhost:9092
 - key.converter=org.apache.kafka.connect.json.JsonConverter
 - value.converter=org.apache.kafka.connect.json.JsonConverter
 - key.converter.schemas.enable=false
 - value.converter.schemas.enable=false
 - offset.storage.file.filename=/tmp/connect.offsets
 - offset.flush.interval.ms=10000
 - plugin.path=/share/java

Quick Start Kafka Connect

Worker Config

- ▶ Note that *plugin.path* can hold a list of paths, where connector implementations are available
- ▶ As we'll use connectors bundled with Kafka, we can set *plugin.path* to *\$KAFKA_HOME/share/java*. Working with Windows, it might be necessary to provide an absolute path here.
- ▶ For the other parameters, we can leave the default values:
 - *bootstrap.servers* contains the addresses of the Kafka brokers
 - *key.converter* and *value.converter* define converter classes, which serialize and deserialize the data as it flows from the source into Kafka and then from Kafka to the sink
 - *key.converter.schemas.enable* and *value.converter.schemas.enable* are converter-specific settings
 - ***offset.storage.file.filename* is the most important setting when running Connect in standalone mode: it defines where Connect should store its offset data**
 - *offset.flush.interval.ms* defines the interval at which the worker tries to commit offsets for tasks

Quick Start Kafka Connect

Kafka Connect in Standalone Mode

- ▶ And with that, we can start our first connector setup:
 - `$KAFKA_HOME/bin/connect-standalone.sh \`
`$KAFKA_HOME/config/connect-standalone.properties \`
`$KAFKA_HOME/config/connect-file-source.properties \`
`$KAFKA_HOME/config/connect-file-sink.properties`
- ▶ First off, we can inspect the content of the topic using the command line:
 - `$KAFKA_HOME/bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic connect-test --from-beginning`

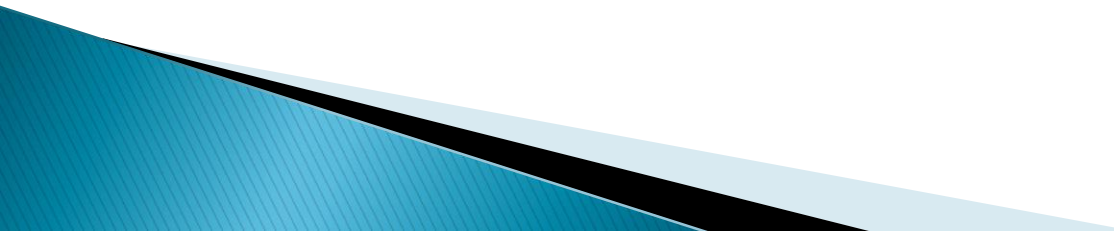
Quick Start Kafka Connect

Kafka Connect in Standalone Mode

- ▶ As we can see, the source connector took the data from the *test.txt* file, transformed it into JSON, and sent it to Kafka:
 - {"schema":{"type":"string","optional":false},"payload":"foo"}
 - {"schema":{"type":"string","optional":false},"payload":"bar"}
- ▶ And, if we have a look at the folder *\$KAFKA_HOME*, we can see that a file *test.sink.txt* was created here:
 - cat *\$KAFKA_HOME/test.sink.txt*
 - foo
 - Bar
- ▶ As the sink connector extracts the value from the *payload* attribute and writes it to the destination file, the data in *test.sink.txt* has the content of the original *test.txt* file.

Quick Start Kafka Connect

Kafka Connect in Standalone Mode

- ▶ Now let's add more lines to *test.txt*.
 - ▶ When we do, we see that the source connector detects these changes automatically.
 - ▶ We only have to make sure to insert a newline at the end, otherwise, the source connector won't consider the last line.
 - ▶ At this point, let's stop the Connect process, as we'll start Connect in *distributed mode* in a few lines.
- 

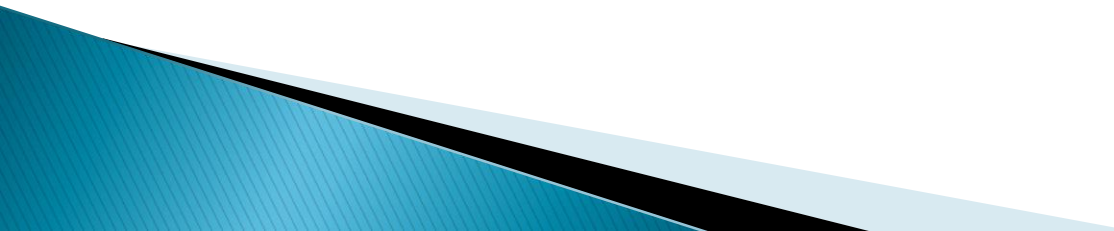
Connect's REST API

- ▶ Until now, we made all configurations by passing property files via the command line.
- ▶ However, as Connect is designed to run as a service, there is also a REST API available.
- ▶ By default, it is available at *http://localhost:8083*. A few endpoints are:
 - *GET /connectors* – returns a list with all connectors in use
 - *GET /connectors/{name}* – returns details about a specific connector
 - *POST /connectors* – creates a new connector; the request body should be a JSON object containing a string name field and an object config field with the connector configuration parameters

Connect's REST API

- *GET /connectors/{name}/status* – returns the current status of the connector – including if it is running, failed or paused – which worker it is assigned to, error information if it has failed, and the state of all its tasks
- *DELETE /connectors/{name}* – deletes a connector, gracefully stopping all tasks and deleting its configuration
- *GET /connector-plugins* – returns a list of connector plugins installed in the Kafka Connect cluster

Kafka Connect in Distributed Mode

- ▶ The standalone mode works perfectly for development and testing, as well as smaller setups.
 - ▶ However, if we want to make full use of the distributed nature of Kafka, we have to launch Connect in distributed mode.
 - ▶ By doing so, connector settings and metadata are stored in Kafka topics instead of the file system.
 - ▶ As a result, the worker nodes are really stateless.
- 

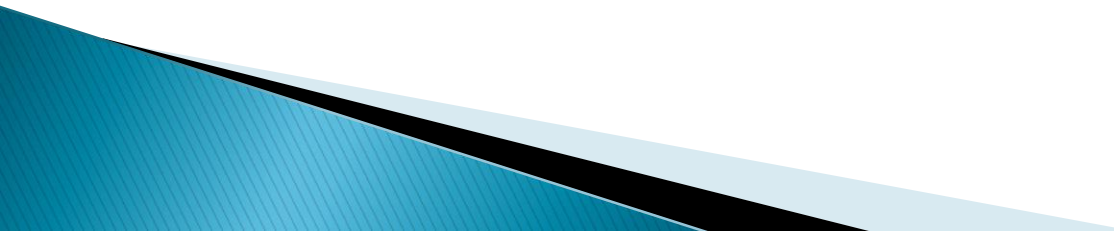
Kafka Connect in Distributed Mode

Starting Connect

- ▶ A reference configuration for distributed mode can be found at `$KAFKA_HOME/config/connect-distributed.properties`.
- ▶ Parameters are mostly the same as for standalone mode. There are only a few differences:
 - ***group.id*** defines the name of the Connect cluster group. The value must be different from any consumer group ID
 - ***offset.storage.topic***, ***config.storage.topic*** and ***status.storage.topic*** define topics for these settings. For each topic, we can also define a replication factor
- ▶ We can start Connect in distributed mode as follows:
 - `$KAFKA_HOME/bin/connect-distributed.sh \`
`$KAFKA_HOME/config/connect-distributed.properties`

Kafka Connect in Distributed Mode

Adding Connectors Using the REST API

- ▶ Now, compared to the standalone startup command, we didn't pass any connector configurations as arguments.
 - ▶ Instead, we have to create the connectors using the REST API.
 - ▶ To set up our example from before, we have to send two POST requests to *http://localhost:8083/connectors* containing the following JSON structs.
- 

Kafka Connect in Distributed Mode

Adding Connectors Using the REST API

- ▶ First, we need to create the body for the source connector POST as a JSON file. Here, we'll call it *connect-file-source.json*:

```
{
  "name": "local-file-source",
  "config": {
    "connector.class": "FileStreamSource",
    "tasks.max": 1,
    "file": "test-distributed.txt",
    "topic": "connect-distributed"
  }
}
```

- ▶ Note how this looks pretty similar to the reference configuration file we used the first time.
- ▶ And then we POST it:
 - `curl -d @"$KAFKA_HOME/connect-file-source.json" \`
 `-H "Content-Type: application/json" \`
 `-X POST http://localhost:8083/connectors`

Kafka Connect in Distributed Mode

Adding Connectors Using the REST API

- ▶ Then, we'll do the same for the sink connector, calling the file *connect-file-sink.json*:

```
{
  "name": "local-file-sink",
  "config": {
    "connector.class": "FileStreamSink",
    "tasks.max": 1,
    "file": "test-distributed.sink.txt",
    "topics": "connect-distributed"
  }
}
```

- ▶ And perform the POST like before:
 - `curl -d @$KAFKA_HOME/connect-file-sink.json \`
`-H "Content-Type: application/json" \`
`-X POST http://localhost:8083/connectors`

Kafka Connect in Distributed Mode

Adding Connectors Using the REST API

- ▶ If needed, we can verify, that this setup is working correctly:
 - `$KAFKA_HOME/bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic connect-distributed --from-beginning`
 - `{"schema":{"type":"string","optional":false},"payload":"foo"}`
 - `{"schema":{"type":"string","optional":false},"payload":"bar"}`
- ▶ And, if we have a look at the folder `$KAFKA_HOME`, we can see that a file `test-distributed.sink.txt` was created here:
 - `cat $KAFKA_HOME/test-distributed.sink.txt`
 - foo
 - Bar
- ▶ After we tested the distributed setup, let's clean up, by removing the two connectors:
 - `curl -X DELETE http://localhost:8083/connectors/local-file-source`
 - `curl -X DELETE http://localhost:8083/connectors/local-file-sink`

Transforming Data

Supported Transformations

- ▶ Transformations enable us to make simple and lightweight modifications to individual messages.
- ▶ Kafka Connect supports the following built-in transformations:
 - *InsertField* – Add a field using either static data or record metadata
 - *ReplaceField* – Filter or rename fields
 - *MaskField* – Replace a field with the valid null value for the type (zero or an empty string, for example)
 - *HoistField* – Wrap the entire event as a single field inside a struct or a map

Transforming Data

Supported Transformations

- ▶ Kafka Connect supports the following built-in transformations:
 - *ExtractField* – Extract a specific field from struct and map and include only this field in the results
 - *SetSchemaMetadata* – Modify the schema name or version
 - *TimestampRouter* – Modify the topic of a record based on original topic and timestamp
 - *RegexRouter* – Modify the topic of a record based on original topic, a replacement string, and a regular expression

Transforming Data

Supported Transformations

- ▶ A transformation is configured using the following parameters:
 - *transforms* – A comma-separated list of aliases for the transformations
 - *transforms.\$alias.type* – Class name for the transformation
 - *transforms.\$alias.\$transformationSpecificConfig* – Configuration for the respective transformation

Transforming Data

Applying a Transformer

- ▶ To test some transformation features, let's set up the following two transformations:
 - First, let's wrap the entire message as a JSON struct
 - After that, let's add a field to that struct
- ▶ Before applying our transformations, we have to configure Connect to use schemaless JSON, by modifying the *connect-distributed.properties*:
 - `key.converter.schemas.enable=false`
 - `value.converter.schemas.enable=false`
- ▶ After that, we have to restart Connect, again in distributed mode:
 - `$KAFKA_HOME/bin/connect-distributed.sh \`
`$KAFKA_HOME/etc/kafka/connect-distributed.properties`

Transforming Data

Applying a Transformer

- ▶ Again, we need to create the body for the source connector POST as a JSON file. Here, we'll call it *connect-file-source-transform.json*.
- ▶ Besides the already known parameters, we add a few lines for the two required transformations:

```
{
  "name": "local-file-source",
  "config": {
    "connector.class": "FileStreamSource",
    "tasks.max": 1,
    "file": "test-transformation.txt",
    "topic": "connect-transformation",
    "transforms": "MakeMap,InsertSource",
    "transforms.MakeMap.type": "org.apache.kafka.connect.transforms.HoistField$Value",
    "transforms.MakeMap.field": "line",
    "transforms.InsertSource.type": "org.apache.kafka.connect.transforms.InsertField$Value",
    "transforms.InsertSource.static.field": "data_source",
    "transforms.InsertSource.static.value": "test-file-source"
  }
}
```

Transforming Data

Applying a Transformer

- ▶ After that, let's perform the POST:
 - `curl -d @$KAFKA_HOME/connect-file-source-transform.json \`
`-H "Content-Type: application/json" \`
`-X POST http://localhost:8083/connectors`
- ▶ Let's write some lines to our *test-transformation.txt*:
 - Foo
 - Bar
- ▶ If we now inspect the *connect-transformation* topic, we should get the following lines:
 - `{"line":"Foo","data_source":"test-file-source"}`
 - `{"line":"Bar","data_source":"test-file-source"}`

Using Ready Connectors

- ▶ After using these simple connectors, let's have a look at more advanced ready-to-use connectors, and how to install them.

Where to Find Connectors

- ▶ Pre-built connectors are available from different sources:
 - A few connectors are bundled with plain Apache Kafka (source and sink for files and console)
 - Some more connectors are bundled with Confluent Platform (ElasticSearch, HDFS, JDBC, and AWS S3)

Using Ready Connectors

- ▶ Also check out [Confluent Hub](#), which is kind of an app store for Kafka connectors. The number of offered connectors is growing continuously:
 - Confluent connectors (developed, tested, documented and are fully supported by Confluent)
 - Certified connectors (implemented by a 3rd party and certified by Confluent)
 - Community-developed and -supported connectors
- ▶ Beyond that, Confluent also provides a [Connectors Page](#), with some connectors which are also available at the Confluent Hub, but also with some more community connectors
- ▶ And finally, there are also vendors, who provide connectors as part of their product. For example, Landoop provides a streaming library called [Lenses](#), which also contains a set of ~25 open source connectors (many of them also cross-listed in other places)

Using Ready Connectors

Installing Connectors Manually

- ▶ We can install the required connectors manually. For that, we have to download and unzip the connector, as well as move the included libs to the folder specified as *plugin.path*.
- ▶ For each connector, the archive should contain two folders that are interesting for us:
- ▶ The *lib* folder contains the connector jar, for example, *kafka-connect-mqtt-1.0.0-preview.jar*, as well as some more jars required by the connector
- ▶ The *etc* folder holds one or more reference config files

Using Ready Connectors

Installing Connectors Manually

- ▶ We have to move the *lib* folder to *\$KAFKA_HOME/share/java*, or whichever path we specified as *plugin.path* in *connect-standalone.properties* and *connect-distributed.properties*. In doing so, it might also make sense to rename the folder to something meaningful.
- ▶ We can use the config files from *etc* either by referencing them while starting in standalone mode, or we can just grab the properties and create a JSON file from them.